

COP4634: Systems & Networks I

Synchronization

3

Background on Process Synchronization

- Concurrent access to shared data may result in data becoming inconsistency.
- Maintaining data consistency while supporting concurrent processing requires additional OS mechanisms
 - shared data may be accessed by multiple processes simultaneously
 - outcome of operation on shared data may depend on the order by which processes execute instructions leading to inconsistent results

L07 Synchronizatio

COB4634

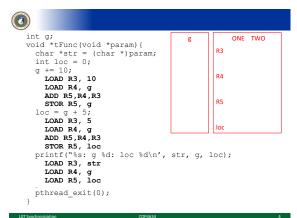


Example: Concurrent Data Access

```
int g;
void *tFunc(void *param) {
    char *str = (char *)param;
    int loc = 0;
    g += 10;
    loc = g + 5;
    printf("%s: g %d: loc %d\n', str, g, loc);
    pthread_exit(0);
}
int main(int argc, char **argv) {
    pthread_t tid1, tid2;
    g = 10;
    pthread_create(&tid1, NULL, tFunc, (void *)"ONE");
    pthread_create(&tid2, NULL, tFunc, (void *)"TWO");
    pthread_join(tid1, NULL);
    return 0;
```

L07 Synchronization

COP4634



Synchronization Concepts

Race Condition: situation in which the final outcome of a result is determined by the order of process execution.

- Different executions produce different results
- Almost always bad
- Scheduling allows order to be imposed

<u>Critical Section</u>: a code section in a process that accesses shared data.

- need to be protected using synchronization mechanisms to ensure data integrity
- while one process enters its critical section, no other process may enter their critical section to manipulate shared data

L07 Synchronizati

COBVESV

5



Solution to Critical Section Problem

- Mutual Exclusion: prohibit other processes from entering their critical section while a process executes its critical section
- Progress: if no process is in their critical section, there must be progress on any process that wants to enter their critical section.
- Bounded waiting: there must be a bound on the number of times other processes are allowed to enter their critical section after a process has made a request to enter its critical section.

L07 Synchronizat

COP4634

-

1	10.07	
18	1	M
ĸ.	С6.	g;)
V	Sec. of	

Critical Sections and Threads

- Threads must proceed through an entry and exit sections.
- Threads must not die or quit inside a critical
- Threads may be context switched inside a critical section.
- A context switch is different from an exit because another thread may not be allowed to enter their critical section.

COP4634



Bad Solution 1

```
int turn;

1 while(1) {
2    ...
3    while (turn != i);
4    CS();
5    turn = j;
6    ...
7 }
```

- •CS Critical Section
- •Mutal exclusion (mutex)
- i pid of currently executing process
- j pid of "other" process

L07 Synchronization

OB4634



Bad Solution 1

```
int turn;

1 while(1) {
2    ...
3    while (turn != i);
4    CS();
5    turn = j;
6    ...
7 }
```

turn PO P1

L07 Synchronizatio

Bad Solution 2

7

General Format

```
while (1) {
...
ENTRY_CODE(); Insures MUTEX
CS();
EXIT_CODE(); Allows next P to enter CS
...
Assume no crashes for now
```

7

Atomic Actions

- Once started, must complete
- Can't be interrupted in the middle
- Completes or doesn't start
- Statements (generally) NOT atomic
- Instructions are atomic

L07 Synchronization

COP4634

3	Synchronization Solutions
	Peterson's and Bakery Algorithm — work but are difficult to generalize to all sorts of computing problems requiring synchronization
	Synchronization Hardware
	atomic instructions that test and set a boolean value at the same time
LO7 Sys	nchronization COP4634 13
7	Problems
•	User-level implementation
	 Can be mistyped by user
	- Can be intentionally skipped
	– Busy wait loop What we want
	Kernel implementation
	– User can use
	– User cannot alter
	Efficient implementation
107.5	200171
LU/Syi	nctronization CDP4634 14
(3)	LOCK
•	Kernel implementation of MUTEX
	Has two functions
	 Acquire() requests lock, only returns when lock is given to process
	 Release() gives lock back to kernel to reallocate to another process
•	Presented as "Class/Object" for clarity
•	Really implemented in C
LO7 Sys	nchronization COP4634 15

```
class LOCK {
  int status = FREE;
}
LOCK::Acquire() {
  Disable_Interrupts();
  while (status == BUSY) {
    Enable_Interrupts();
    Disable_Interrupts();
  }
  status = BUSY;
  Enable_Interrupts();
}
```

class LOCK {
 int status = FREE;
}
LOCK::Acquire() {
 ...
}
LOCK::Release() {
 Disable_Interrupts();
 status = FREE;
 Enable_Interrupts();

3

Why not Perfect Solution?

- · Busy wait loop still present
 - Eats CPU time while waiting
- · High priority thread waiting?
 - Holder (low priority) never runs
- · Waiting thread could be "infinitely unlucky"
 - How could this be?

L07 Synchronization

COP4634

```
class LOCK {
  int status = FREE;
}
LOCK::Acquire() {
  Disable_Interrupts();
  while (status == BUSY) {
    Enable_Interrupts();
    Disable_Interrupts();
}
status = BUSY;
Enable_Interrupts();
}

PO
interrupted
here

P1 Release(),
Acquire(),
status is FREE
```

Corrections

- Add "Waiting List for Lock" (queue)
- Need enqueue () and dequeue () for list
- Must be able to put "self" on list
- Idea
 - If lock is BUSY,
 - put self on waiting list
 - go to sleep
 - When releasing lock,
 - wake a sleeping thread

LO7 Synchronizatio

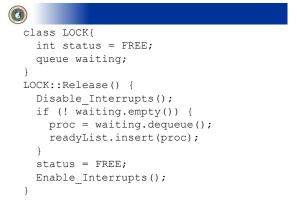
COP4634

(3)

```
class LOCK{
  int status = FREE;
  queue waiting;
}
LOCK::Acquire() {
  Disable_Interrupts();
  while (status != FREE) {
    waiting.enqueue(self);
    Enable_Interrupts();
    block();
    Disable_Interrupts();
}
status = BUSY;
Enable_Interrupts();
}
```

L07 Synchronizatio

COP4634



(3)

Notes

- block () is internal kernel call (kernel mode)
- Interrupts must be enabled before calling block()
- · Similar to context switch
 - Calling process taken off CPU
 - Calling process stored in PCB
 - Calling process **NOT** put on ready list
- Better be on some other list before calling!

L07 Synchronizatio

COP4634

73



LOCKs in Pthreads

```
#include <pthread.h>
pthread_mutex_t lock;
pthread_mutexattr_t mattr;
pthread_mutexattr_init( &mattr );
pthread_mutex_init( &lock, &mattr );
pthread_mutex_lock( &lock );
pthread_mutex_unlock( &lock );
pthread_mutex_destroy( &lock );
```

L07 Synchronization

COP4634



Hardware Help

- Read/Modify/Write instruction
 - test-n-set instruction (atomic instruction)
- · Functional Definition (really single instruction)

Example:
1) inputVal = 0;
2) returnVal = testNset(inputVal);

Q1: What is the value of *inputVal* and *returnVal* after the execution of *testNset()*?
A1: inputVal = 1 and returnVal = 0.

Q2: What changes if *testNset()* is called again with the same value for *inputVal*?

L07 Synchronizati

COP4634

2



Test-N-Set

- · Every modern CPU has equivalent
 - exchange
 - compare-and-swap

```
LOCK::Acquire() {
    // spins a process as long as
    // status is true
    while (testNset(status) == 1);
}
LOCK::Release() {
    status = 0;
```

L07 Synchronization

COP4634

26



Limitations

- · Locks are good for MUTEX only
- General concept is needed to solve synchronization problems:



Semaphores

- Synchronization concept first proposed by Edsger Dijkstra.
- Semaphores are objects maintained by the OS.
- Each semaphore has a value and two atomic operations:
 - wait
 - signal

L07 Synchronization

COP4634

Other Names & Original Definition

```
"test"
Proberen() { P() Wait() sem_wait()
  while (count == 0);
  count--;
}

"increment"
Verhogen() { V() Signal() sem_post()
  count++;
}
```

Extended Definition

```
int sem_wait( sem ) {
  Disable_Interrupts();
  sem.count--;
  if (sem.count < 0) {
    waitlist.enqueue(self);
    Enable_Interrupts();
    block();
    Disable_Interrupts();
}
Enable_Interrupts();
return 0;
}</pre>
```

Extended Definition

```
int sem_post( sem ) {
   Disable_Interrupts();
   sem.count++;
   if (sem.count <= 0) {
      proc = waitlist.dequeue();
      readyList.insert(proc);
   }
   Enable_Interrupts();
   return 0;
}</pre>
```

L07 Synchronization

COP4634

1	1.0	
1:0	1	76.
170	r.	188
V		20

Monitors for Synchronization

- · Monitors are a high-level synchronization construct.
- Monitors allow safe sharing of abstract data types among concurrent processes.
- Monitors provide synchronized methods that can only be entered by one process or thread at a time.
- Monitors provide condition variables with two operations
 - wait: puts process or thread asleep
 - signal: triggers a waiting process or thread to resume operation <u>exactly at the position</u> where process or thread called wait

L07 Synchroniza

COP4634

31



Condition Variables Implementation (pseudo code)

```
Wait( lock ) {
    // put one thread asleep
    lock.Acquire();
}
Signal( lock ) {
    // wake one sleeping thread
    lock.release();
}
Broadcast( lock ) {
    // wake-up all n threads
    repeat n times
    lock.release();
```

L07 Synchronization

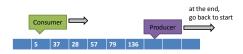
COB4634

22



Producer/Consumer Synchronization Problem

- Producer makes widgets, Consumer eats widgets
- Storage facility (limited size) for widgets
 - Storage is a circular data structure
- Producer must stop generating data when buffer is full.
- Consumer must stop reading data when no data are available.
- Examples: Networks, Disk I/O, etc.



L07 Synchronizati

COP4634

```
sem_t fullB;
sem_t emptyB;
pthread_mutex_t mutex;
void * producer(void *id) {
  int myID = (int)id;
  while(1) {
    sem_wait( &emptyB );
    pthread_mutex_lock( &mutex );
    addWidgetToBuffer();
    pthread_mutex_unlock( &mutex );
    sem_post( &fullB );
}
Q: What is the initial value of emptyB and fullB?
```

(3)

```
sem_t fullB;
sem_t emptyB;
pthread_mutex_t mutex;
void * consumer(void *id) {
  int myID = (int)id;
  while(1) {
    sem_wait( &fullB );
    pthread_mutex_lock( &mutex );
    eatWidgetFromBuffer();
    pthread_mutex_unlock( &mutex );
    sem_post( &emptyB );
}
A: Initially, the value emptyB must be the size of the buffer
```

(7)

File Print

```
sem_wait(&printer); sem_wait(&file);
sem_wait(&file); sem_wait(&printer);
printFile(); printFile();
sem_post(&file); sem_post(&printer);
sem_post(&printer); sem_post(&file);
```

Deadlock is possible

and fullB must be 0.

•Be careful with order

L07 Synchronization

5004534

1	4.0	
68	1	ħ
и.	r.	189
V.	1	

Sleeping Barber

- Small town
- One barber shop w/ one barber
- · Barber always at work
- Sleeps when possible
- 1st customer in shop wakes barber
- · Other customers wait
- All customers done, barber goes back to sleep

3

Sleeping Barber

```
Customer() {
                      mutex.Acquire();
Barber() {
                      if (numWaiting < CHAIRS) {</pre>
while (1) \{
                      numWaiting++;
 barber.Wait();
                       mutex.Release();
 mutex.Acquire();
 numWaiting--;
                       barber.Signal();
 mutex.Release();
                       done.Wait();
 done.Signal();
                      } else
                        mutex.Release();
                      leave();
```

Plato

Descartes

Voltaire

Socrates

(3)

Dining Philosophers Problem Statement

- N philosophers at round table
- · Bowl of rice in middle of table
- Fork between each pair of philosophers
- Need 2 forks to eat
- think, get hungry, grab forks, eat, ...
- · Avoid deadlock
- · Avoid starvation
- · Let as many eat as possible

L07 Synchroniz

COP4634

*Philosophers are identified by number *Prorks are non-preemptive resources.

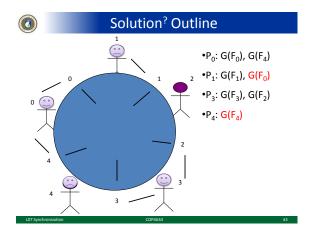
0

How about this?

```
void *philosopherThread(void *id) {
  int myID = (int)id;
  while(1) {
    think();
    grab(myID - 1);
    grab(myID);
    eat();
    replace(myID - 1);
    replace(myID);
}
```

L07 Synchronization

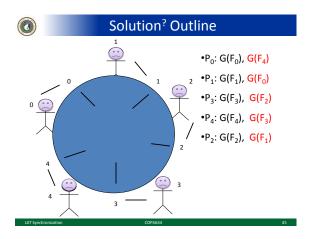
COP4634



What if...

- One philosopher's stomach growls
- All philosophers realize they are hungry
- All philosophers grab left fork
- All philosophers sleep waiting for right fork
- Deadlock!
- How can we fix the problem?







Another Try

```
void *philosopherThread(void *id) {
  int myID = (int)id;
  while(1) {
    think();
    pthread_mutex_lock( &mutex );
    grab(myID - 1);
    grab(myID);
    eat();
    replace(myID - 1);
    replace(myID);
    pthread_mutex_unlock( &mutex );
}
```

Solution? Outline

•P₀: G(F₀), G(F₄)

•P₁: G(F₁)

•P₃: G(F₃)

•P₄: G(F₄)



Did it Work?

- Do we avoid deadlock?
- How many can eat at a time?
- Do you think this is a good solution?
- How can we improve it?
 - Analyze the problem
 - What is really causing deadlock
 - How can we avoid it?
 - Can we make our solution a *little* better?

L07 Synchronizati

COP4634



Shared Queue

```
LOCK lock = FREE;
AddToQueue( item ) {
  lock.Acquire();
  list.append(item);
  lock.Release();
}
RemoveFromQueue() {
  lock.Acquire();
  if(! list.isEmpty())
    item = list.remove();
  lock.Release();
  return item;
}
```

L07 Synchronization

COP4634



Shared Queue

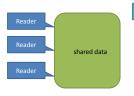
```
RemoveFromQueue() {
  lock.Acquire();
  if( list.isEmpty() )
     lock.Release();
     Sleep();
     lock.Acquire();
  item = list.remove();
  lock.Release();
  return item;
}
```

L07 Synchronizatio

COP4634

Multiple Reader, Single Writer

- Multiple Reader processes may access simultaneously shared data, provided Writer process has no access to it.
- Writer process may not access shared data unless no Reader process accesses shared data.



"Writer must wait until all Readers are done reading data."

L07 Synchronizatio

COP4634

(7)

Example - DB Access

```
int activeReaders = 0;
int activeWriters = 0;
int waitingReaders = 0;
int waitingWriters = 0;
Condition OKtoRead;
Condition OKtoWrite;
Lock lock = FREE;
```

L07 Synchronizatio

COP4634

52



Example – DB Access

```
ReadStart() {
    lock.Acquire();
    while( (activeWriters + waitingWriters) > 0 ) {
        waitingReaders++;
        OKtoRead.Wait(lock);
        waitingReaders--;
    }
    activeReaders++;
    lock.Release();
}

ReadStop() {
    lock.Acquire();
    activeReaders--;
    if ( (activeReaders == 0) && (waitingWriters > 0) )
        OKtoWrite.Signal(lock);
    lock.Release();
}
```

L07 Synchronizatio

COP4634

```
3
```

Example – DB Access

```
WriteStart() {
   lock.Acquire();
   while( (activeWriters + activeReaders) > 0 ) {
      waitingWriters++;
      OKtoWrite.Wait(lock);
      waitingWriters--;
   }
   activeWriters++;
   lock.Release();
}
```

L07 Synchronizati

COP4634

Example – DB Access

```
WriteStop() {
   lock.Acquire();
   activeWriters--;
   if ( waitingWriters > 0 )
      OKtoWrite.Signal(lock);
   else if ( waitingReaders > 0 )
      OKtoRead.Broadcast(lock);
   lock.Release();
}
```

LO7 Synchronizatio

COP4634

Usage

LO7 Synchronizatio

COP4634

9/18/2013

3	Com	nparison
	naphore - Has a value	Condition Variable - Doesn't have a value
-	- Implies a history	 Implies no memory
-	 wait() may sleep or may proceed depending 	wait() always sleeps
-	on value - signal() will	<pre>- signal() will wake a</pre>
	increment a value and wake a sleeper	sleeper
-	<pre>- signal() before wait() will effect result of wait</pre>	<pre>- signal() before wait() has no effect</pre>